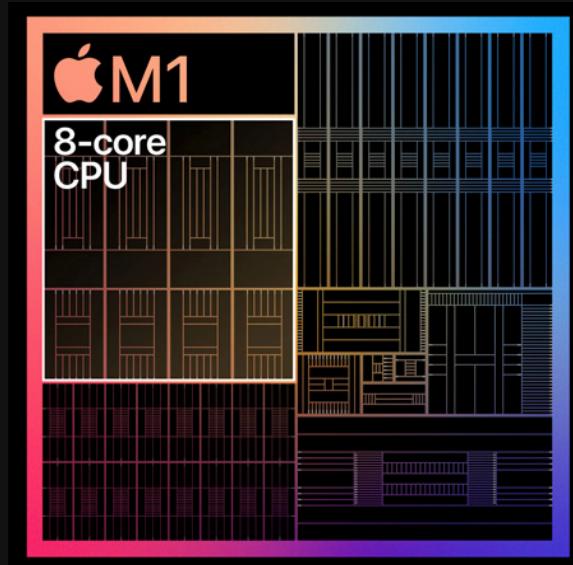


# Introduction

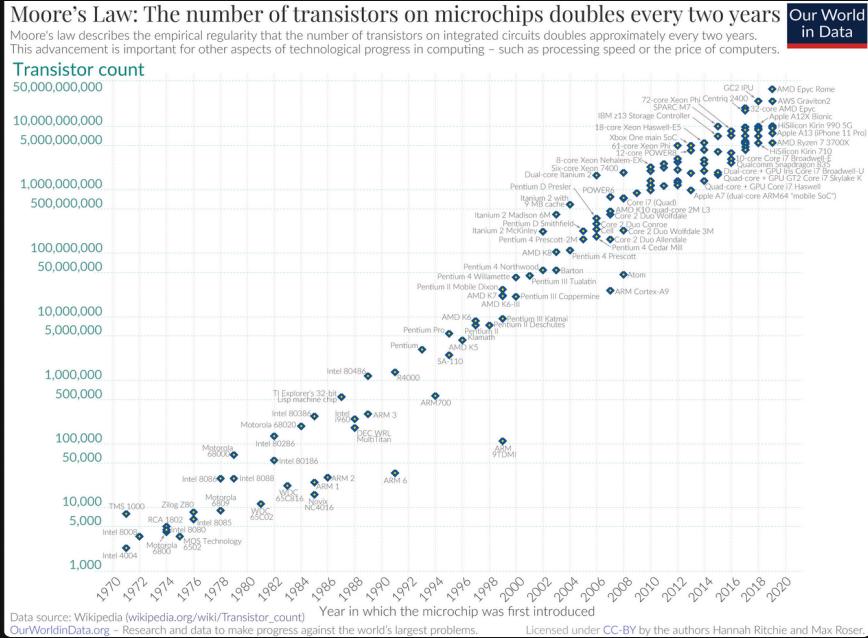
# Introduction

Parallel computing is an important aspect of modern computing systems.

And while you *may* not be writing parallel code ourselves, with the tools we use every day automatically leveraging parallelism, it is important to understand the concepts and principles behind parallel computing.



# Single processor performance



- 1986 - 2003
- 2003
- 2015 - 2017

# Change in processor design

By 2005, most of the major manufacturers of microprocessors had decided that the road to rapidly increasing performance lay in the direction of **parallelism**.

Intel Pentium D and AMD Athlon 64 x 2

For software developers:

- simply adding more processors will not improve the performance of a programs




Why care about parallel computing?

To increase performance

# limitations of single processor performance

- As the size of transistors decreases, their speed can be increased, and the overall speed of the integrated circuit can be increased.
- As the speed of transistors increases, their power consumption also increases.
- Most of this power is dissipated as heat, and air-cooled integrated circuits reached the limits of their ability to dissipate heat



Therefore the only valid path for improved performance is **not** relying on ever-faster single processors

# why write parallel programs

- Why do we have to actually write parallel programs?

Most programs that have been written for conventional, single-core systems **cannot** use the presence of multiple cores

While we can run multiple instances of a program on a multicore system, that's not very useful

- We need to either rewrite our serial programs or write translation programs

## The problem with translation

Researchers have had limited success writing programs that convert serial programs

An efficient parallel implementation of a serial program may not be obtained by finding efficient parallelizations of *each of its steps.*

The best parallelization tends to be from devising an entirely new algorithm

An example: sum

# Problem

Suppose that we need to compute n values and add them together

```
1 [1, 4, 3, 9, 2, 8,  
2 5, 1, 1, 6, 2, 7,  
3 2, 5, 0, 4, 1, 8,  
4 6, 5, 1, 2, 3, 9, ]
```

## An example implementation

```
1 sum = 0;
2 for ( i = 0; i < n; i ++ ) {
3     x = Compute_next_value( ... );
4     sum += x;
5 }
```

# With multiple cores

Assume eight cores

```
1 my_sum = 0;
2 my_first_i = ... ;
3 my_last_i = ... ;
4
5 for(my_i = my_first_i; my_i < my_last_i; my_i++) {
6     my_x = Compute_next_value( ... );
7     my_sum += my_x;
8 }
```

With 24 values and 8 cores, how much should each core compute?

# Global Sum

```
1 if (Im the master core) {  
2     sum = my_sum;  
3     for each core other than myself {  
4         recieve value from core;  
5         sum += value;  
6     }  
7 } else {  
8     send my my_sum to master core;  
9 }
```

# How would you improve this program

hint: how would you lower the amount of time spent in the main core

# How to write parallel programs

# How to write parallel programs

There are a number of possible answers to this question, but most of them depend on the basic idea of **partitioning** the work to be done among the cores

There are some common approaches for this, but before that

## Example

- Suppose that I have to teach a section of this subject
- Say that I have **100** students
- At the end of the semester, I have to grade their final exams
- So I recruit **4** student assistants to help me check
- the final exam consists of **five** open ended questions with a rubric

How would I use the five of us to grade the exams as quickly as possible?

# Task Parallelism

I can give each of the 4 assistants a **task**

- SA 1 grades only question 1,
- SA 2 grades only question 2,
- and so on.

The total task is *partitioned* into five **different** tasks, each of which is done by one person (core)

This is called **task parallelism**

# Data Parallelism

Alternatively, I can divide the one hundred exams into **five piles** of twenty exams each,

- SA 1 grades all the papers in the first pile,
- SA 2 grades all the papers in the second pile,
- and so on.

The total task is *partitioned* is not partitioned, but the **data** to be processed is partitioned into five **similar** tasks, each of which is done by one person

Give me pros and cons of each scenario

Back to the sum example

Is it task or data parallelism?

# Coordination

# Coordination in Parallel Programs

*if the tasks can be done independently*

Writing parallel programs can be easy

When cores can work independently, it works the same way as serial programs

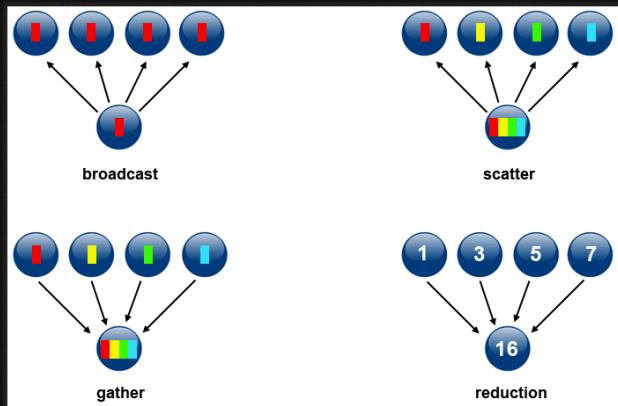
If not, there needs to be **coordination** between tasks

# Coordination in Parallel Programs

The more common ways for tasks to be coordinated are:


## Communication

- One or more cores **send** their current partial sums to another core
- **Sending** and **receiving** data



## Load balancing

- we want the amount of time taken by each core to be roughly **the same**
- if one core has to do more work, we **waste** computing resources



# Synchronization

Instead of computing the values to be added, the values are human inputs.

Say, `x` is an array that is read in by the master core:

```
1 if ( I'm the master core )
2     for (my_i = 0; my_i < n; my_i++)
3         scanf("%d", &x[my_i]);
```

In most systems the cores are **not automatically synchronized**.

- Each core works at its own pace.

If we assume that the core is running the following code:

```
1 for (my_i = my_first_i; my_i < my_last_i; my_i++)
2     my_sum += x[my_i];
```

Give me one runtime error

# Synchronization

In this case, the cores need to **wait** before starting execution of the code:

```
1  for (my_i = my_first_i; my_i < my_last_i; my_i++)
2      my_sum += x[my_i];
```

We need to add in a point of **synchronization** between the initialization of `x` and the computation of the `partial sums`:

```
1  Sync_cores();
```

The idea here is that each core will wait **in the function** `Sync_cores` until **all** the cores have entered the function

So given, `input`, `compute partial`, `sync`, and `compute global`, what is the correct order of execution?

# Quiz in NEO

15-minute study break