

Parallel hardware

Flynn's Taxonomy

In 1966, Michael J. Flynn proposed a classification system for computer architectures based on

Number of concurrent instruction and data streams they support.

The four categories are:


1. **SISD** (Single Instruction, Single Data)
2. **SIMD** (Single Instruction, Multiple Data)
3. **MISD** (Multiple Instruction, Single Data)
4. **MIMD** (Multiple Instruction, Multiple Data)

Which one of these is the classic von Neumann architecture?

Almost every single parallel system is *multiple data*, and the main difference is if they support **one** instruction stream, or **multiple** instruction streams.

SIMD

Applies the **same** instruction to **multiple** data points.



Say we want to add the contents of two arrays, assuming we have an **SIMD** system with n datapaths,

we could load `one[i]` and `two[i]` into the i th datapath (core), then store it at `one[i]`

SIMD

What if we only want to add `two[i]` to `one[i]` if `two[i]` is positive?

```
1  for (i = 0; i < n; i++) {  
2      if (y[i] > 0.0) {  
3          x[i] = x[i] + y[i];  
4      }  
5  }
```

If `y[i]` isn't positive, the datapath holding `y[i]` is idle for that instruction.

Dividing data and having the processors all apply **the same instructions** to their subsets of the data is called what kind of parallelism?

SIMD parallelism can be very efficient on large data parallel problems, but higher control flow can lead to inefficiencies.

GPUs

The graphics pipeline is uniquely suited to an SIMD architecture

In a typical shader function, the **same** operations are applied to **many pixels** or vertices

```
1 uniform vec2 player_position;
2 uniform vec4 color : hint_color = vec4(0.305, 0.298, 0.341,1);
3
4 uniform float MULTIPLIER = 0.56;
5 uniform float SCALE = 0.5;
6 uniform float SOFTNESS = 0.45;
7
8 void fragment(){
9     float val = distance(vec2(UV.x , UV.y * MULTIPLIER),
10                           vec2(player_position.x , player_position.y * MULTIPLIER));
11    val = val / SCALE;
12    float vignette = smoothstep(0.2, SOFTNESS, val);
13    COLOR = vec4(color.rgb, vignette);
14 }
```

How many pixels are there in a 1080 by 1920 display?

GPUs and SIMD

Processing a single image can require very large amounts of data, hundreds of megabytes of data is *normal*

Because SIMD systems **don't** need a full control unit for each datapath, they can have **thousands** of datapaths

The rtx 4090 has **16,384** CUDA cores*




The AMD threadripper has **64**

*Note that these are not equivalent



MIMD


Applies **different** instructions to **different** data points.



Unlike SIMD systems, MIMD systems are usually **asynchronous** and requires **synchronization** between processors.

There are 2 main types of MIMD systems based on how their **memory is arranged**

Shared Memory




Where each core can **share** access to the computer's memory. This style is the most common in **multi-core** processors

Access to the main memory is either **direct**, where all cores can access the main memory.

Or each core can access a **block** of local memory, with access to other core's memory


Shared Memory

That first type, where all cores can access the main memory directly, is called **Uniform Memory Access** (UMA)




These types are usually easier to program, since all cores access the memory at the same latency

The second type, where each core has a block of memory, is called **Non-Uniform Memory Access** (NUMA)




But in a NUMA system, access time is **non-uniform**, which means that some memory accesses are faster than others

Distributed Memory



Where each core has its own **private** memory. This style is the most common in **computer clusters**

Usually used in multi-PC systems connected over a **network**, like Ethernet.



One example of distributed memory systems are **botnets**, where many computers are connected over the internet to perform tasks

Interconnection networks

All parallel systems need a path to communicate

in shared memory systems

In shared memory systems, the main interconnect is the **memory bus**

They are low cost and flexible, but might become slower as traffic increases.



With their main downside being the limit to **distance**, since all cores need to be close enough to the memory bus

in distributed memory systems

In distributed memory systems, the interconnect is usually a **network**



But they're usually divided into two categories:

1. Direct interconnects, where each node is **directly** connected to other nodes (ring, mesh)
2. Indirect interconnects, where nodes are connected through **switches** (star, fat-tree)

Shared vs Distributed

Shared memory systems are easier to program, **faster**, and have lower latency, and

Distributed memory systems are often better for problems require **large amounts of memory**

Directly

A larger bus in a shared memory system is more **expensive**,

quickly becomes a bottleneck for scalability

A larger bus in a distributed memory system is **cheaper**,

but requires more complex routing algorithms